

# Untitled

March 17, 2023

Import necessary libraries: The first few lines import the necessary libraries for loading and visualizing the dataset, training and evaluating the model, and creating the report.

```
[2]: import pandas as pd
import matplotlib.pyplot as plt
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score
from sklearn.model_selection import train_test_split
```

Load the dataset: The `pd.read_csv()` function is used to load the dataset from a CSV file and create a `pandas` dataframe.

```
[8]: # Load the dataset from a CSV file
df = pd.read_csv("dataset.csv")
df.head
```

```
[8]: <bound method NDFrame.head of
feature_1 feature_2 feature_3 feature_4
feature_5 feature_6 \
0 -1.085631 0.997345 0.282978 -1.506295 -0.578600 1.651437
1 -0.678886 -0.094709 1.491390 -0.638902 -0.443982 -0.434351
2 0.737369 1.490732 -0.935834 1.175829 -1.253881 -0.637752
3 -0.255619 -2.798589 -1.771533 -0.699877 0.927462 -0.173636
4 -0.805367 -1.727669 -0.390900 0.573806 0.338589 -0.011830
.. ...
995 -0.462125 0.357491 0.495822 1.286204 -0.695551 -0.065759
996 0.797611 -0.354865 -1.293321 -0.101811 -0.595784 0.117276
997 -0.585542 0.086091 -0.101158 -0.418401 0.490092 0.703430
998 0.293470 0.652308 0.567850 1.478910 -0.400071 -1.667558
999 -1.060740 -1.285079 1.071255 1.910933 0.649171 -0.597918

feature_7 feature_8 feature_9 feature_10 target
0 -2.426679 -0.428913 1.265936 -0.866740 1.0
1 2.205930 2.186786 1.004054 0.386186 1.0
2 0.907105 -1.428681 -0.140069 -0.861755 0.0
3 0.002846 0.688223 -0.879536 0.283627 1.0
4 2.392365 0.412912 0.978736 2.238143 0.0
.. ...
995 0.289966 -0.760412 1.437451 -0.335668 1.0
996 -0.114699 0.745163 -0.635988 0.848133 0.0
```

```

997    1.397362  -1.419400   2.826198    0.793412    1.0
998   -1.350780  -0.112624   0.539086   -0.537706    0.0
999   -1.397932  -0.472266   0.581964    0.970613    1.0

```

```
[1000 rows x 11 columns]>
```

```
[4]: # Data exploration
print(df.describe())
```

	feature_1	feature_2	feature_3	feature_4	feature_5 \
count	1000.000000	1000.000000	1000.000000	1000.000000	1000.000000
mean	-0.005297	0.012115	0.017455	-0.011119	-0.008053
std	0.986820	0.998430	0.990776	0.978192	1.003635
min	-3.167055	-3.685499	-3.066988	-3.581474	-3.587494
25%	-0.693351	-0.684767	-0.685417	-0.702782	-0.670098
50%	0.007288	-0.005697	-0.004751	0.052642	-0.025822
75%	0.641805	0.671780	0.713492	0.660364	0.672929
max	3.050755	3.571579	3.386115	2.789487	3.558981


  

	feature_6	feature_7	feature_8	feature_9	feature_10 \
count	1000.000000	1000.000000	1000.000000	1000.000000	1000.000000
mean	0.044696	0.054121	-0.056954	0.028626	0.021528
std	0.986880	1.016036	1.015486	1.023492	0.980629
min	-3.231055	-3.570243	-3.316010	-3.801378	-3.114263
25%	-0.598035	-0.636259	-0.776020	-0.637296	-0.647347
50%	0.082642	0.060037	-0.064200	0.044113	0.035680
75%	0.690880	0.741102	0.618810	0.714136	0.707607
max	3.569280	2.996198	4.068097	3.311735	2.958625

	target
count	1000.000000
mean	0.515000
std	0.500025
min	0.000000
25%	0.000000
50%	1.000000
75%	1.000000
max	1.000000

Visualize the dataset: The `plt.scatter()` function is used to create a scatter plot of the dataset, with different colors for each target class. The resulting plot is saved to a PNG file using `plt.savefig()`



Split the data into training and testing sets: The `train_test_split()` function from scikit-learn is used to split the dataset into training and testing sets.

```
[5]: # Split the data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(df.iloc[:, :-1], df.iloc[:, -1],
    ↪ test_size=0.2, random_state=123)
```

Train a logistic regression model: The `LogisticRegression()` function is used to create a logistic regression model, which is then trained on the training set using the `fit()` method.

```
[6]: # Train a logistic regression model
model = LogisticRegression(random_state=123)
model.fit(X_train, y_train)
```

```
[6]: LogisticRegression(random_state=123)
```


Evaluate the model on the testing set: The trained model is used to make predictions on the testing set using the `predict()` method, and the accuracy of the predictions is calculated using `accuracy_score()` from scikit-learn.

```
[7]: # Evaluate the model on the testing set
y_pred = model.predict(X_test)
accuracy = accuracy_score(y_test, y_pred)
# Print the model accuracy
```

```
print(f"Model accuracy: {accuracy:.2f}")
```

Model accuracy: 0.50

Visualize the model predictions on the testing set: The `plt.scatter()` function is used again to create a scatter plot of the testing set predictions, with different colors for each predicted class. The resulting plot is saved to a PNG file using `plt.savefig()`.



[ ]: